


# Bagging

Bagging, also known as Bootstrap aggregating, is an ensemble learning technique that helps to improve the performance and accuracy of machine learning algorithms. It is used to deal with bias-variance trade-offs and reduces the variance of a prediction model. Bagging avoids overfitting of data and is used for both regression and classification models, specifically for decision tree algorithms.



## Steps to Perform Bagging

- Consider there are  $n$  observations and  $m$  features in the training set. You need to select a random sample from the training dataset without replacement
- A subset of  $m$  features is chosen randomly to create a model using sample observations
- The feature offering the best split out of the lot is used to split the nodes
- The tree is grown, so you have the best root nodes
- The above steps are repeated  $n$  times. It aggregates the output of individual decision trees to give the best prediction

### Advantages of Bagging in Machine Learning

- Bagging minimizes the overfitting of data
- It improves the model's accuracy
- It deals with higher dimensional data efficiently

## Steps involved

### Data Injection

- Data Profiling
- Basic Operations
- Data Cleaning
- Analysis of features and Statistical Analysis

### EDA

- Univariate Analysis
- Bivariate Analysis
- Multivariate Analysis
- Pre-processing
- Handling DUPLICATE values
- Handling null values

### Mapping

- Feature Encoding
- Splitting of categorical and numerical variable
- Train-Test split

### Model Creation

- Decision Tree Classifier
- HyperParameter Tuning : Decision Tree Classifier
- Bagging Classifier
- Hyperparameter tuning : Bagging Classifier
- Random Forest Classifier
- Hyperparameter tuning : Random Forest Classifier
- Extra Trees Classifier
- HyperParameter Tuning : Extra Tree Classifier
- Voting Classifier
- hard\_voting
- soft\_voting

Evaluation

- Accuracy Score
- Roc-auc score
- Precision
- Recall
- F1\_Score

Attribute Information:

Listing of attributes:

50K, <=50K.

- age: continuous
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, 3. State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, 6. Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

In [30]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import scipy.stats as stats
import warnings
warnings.filterwarnings('ignore')

# Visualization
import matplotlib.pyplot as plt
%matplotlib inline
import missingno
import seaborn as sns
from pandas.plotting import scatter_matrix
from mpl_toolkits.mplot3d import Axes3D
import math

from sklearn.preprocessing import OneHotEncoder, LabelEncoder, label_binarize
from sklearn.model_selection import train_test_split
#Model
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
# Evaluation Metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import roc_curve
# Plots
from sklearn import tree
```

Data Ingestion

In [2]:

```
df = pd.read_excel('Income_data.xlsx')
len(df)
```

Out[2]:

48842

Numerical Columns

```
In [3]:
cat_col=[fea for fea in df.columns if df[fea].dtype == 'O']
cat_col

Out[3]:
['worktype',
 'education',
 'maritial_status',
 'occupation',
 'relationship',
 'race',
 'gender',
 'native_country',
 'salary']
```

Categorical Columns

```
In [4]:
num_col=[fea for fea in df.columns if df[fea].dtype != 'O']
num_col

Out[4]:
['age',
 'fnlwgt',
 'education_num',
 'capital_gain',
 'capital_loss',
 'hours_per_week']
```

Univariate Analysis

In [5]:

```
# Let's plot the distribution of each feature
def plot_distribution(df, cols=5, width=20, height=15, hspace=0.2, wspace=0.5):
    plt.style.use('seaborn-whitegrid')
    fig = plt.figure(figsize=(width,height))
    fig.subplots_adjust(left=None, bottom=None, right=None, top=None, wspace=wspace, hspace=hspace)
    rows = math.ceil(float(df.shape[1]) / cols)
    for i, column in enumerate(df.columns):
        ax = fig.add_subplot(rows, cols, i + 1)
        ax.set_title(column)
        if df.dtypes[column] == np.object:
            g = sns.countplot(y=column, data=df)
            substrings = [s.get_text()[:18] for s in g.get_yticklabels()]
            g.set(yticklabels=substrings)
            plt.xticks(rotation=25)
        else:
            g = sns.distplot(df[column])
            plt.xticks(rotation=25)
```

```
plot_distribution(df, cols=3, width=20, height=20, hspace=0.45, wspace=0.5)
```




Bivariate Analysis

In [6]:

```
# Plot a count of the categories from each categorical feature split by our prediction class: salary - predclass.
def plot_bivariate_bar(dataset, hue, cols=5, width=20, height=15, hspace=0.2, wspace=0.5):
    dataset = dataset.select_dtypes(include=[np.object])
    plt.style.use('seaborn-whitegrid')
    fig = plt.figure(figsize=(width,height))
    fig.subplots_adjust(left=None, bottom=None, right=None, top=None, wspace=wspace, hspace=hspace)
    rows = math.ceil(float(dataset.shape[1]) / cols)
    for i, column in enumerate(dataset.columns):
        ax = fig.add_subplot(rows, cols, i + 1)
        ax.set_title(column)
        if dataset.dtypes[column] == np.object:
            g = sns.countplot(y=column, hue=hue, data=dataset)
            substrings = [s.get_text()[:10] for s in g.get_yticklabels()]
            g.set(yticklabels=substrings)

plot_bivariate_bar(df, hue='salary', cols=1, width=20, height=40, hspace=0.4, wspace=0.5)
```



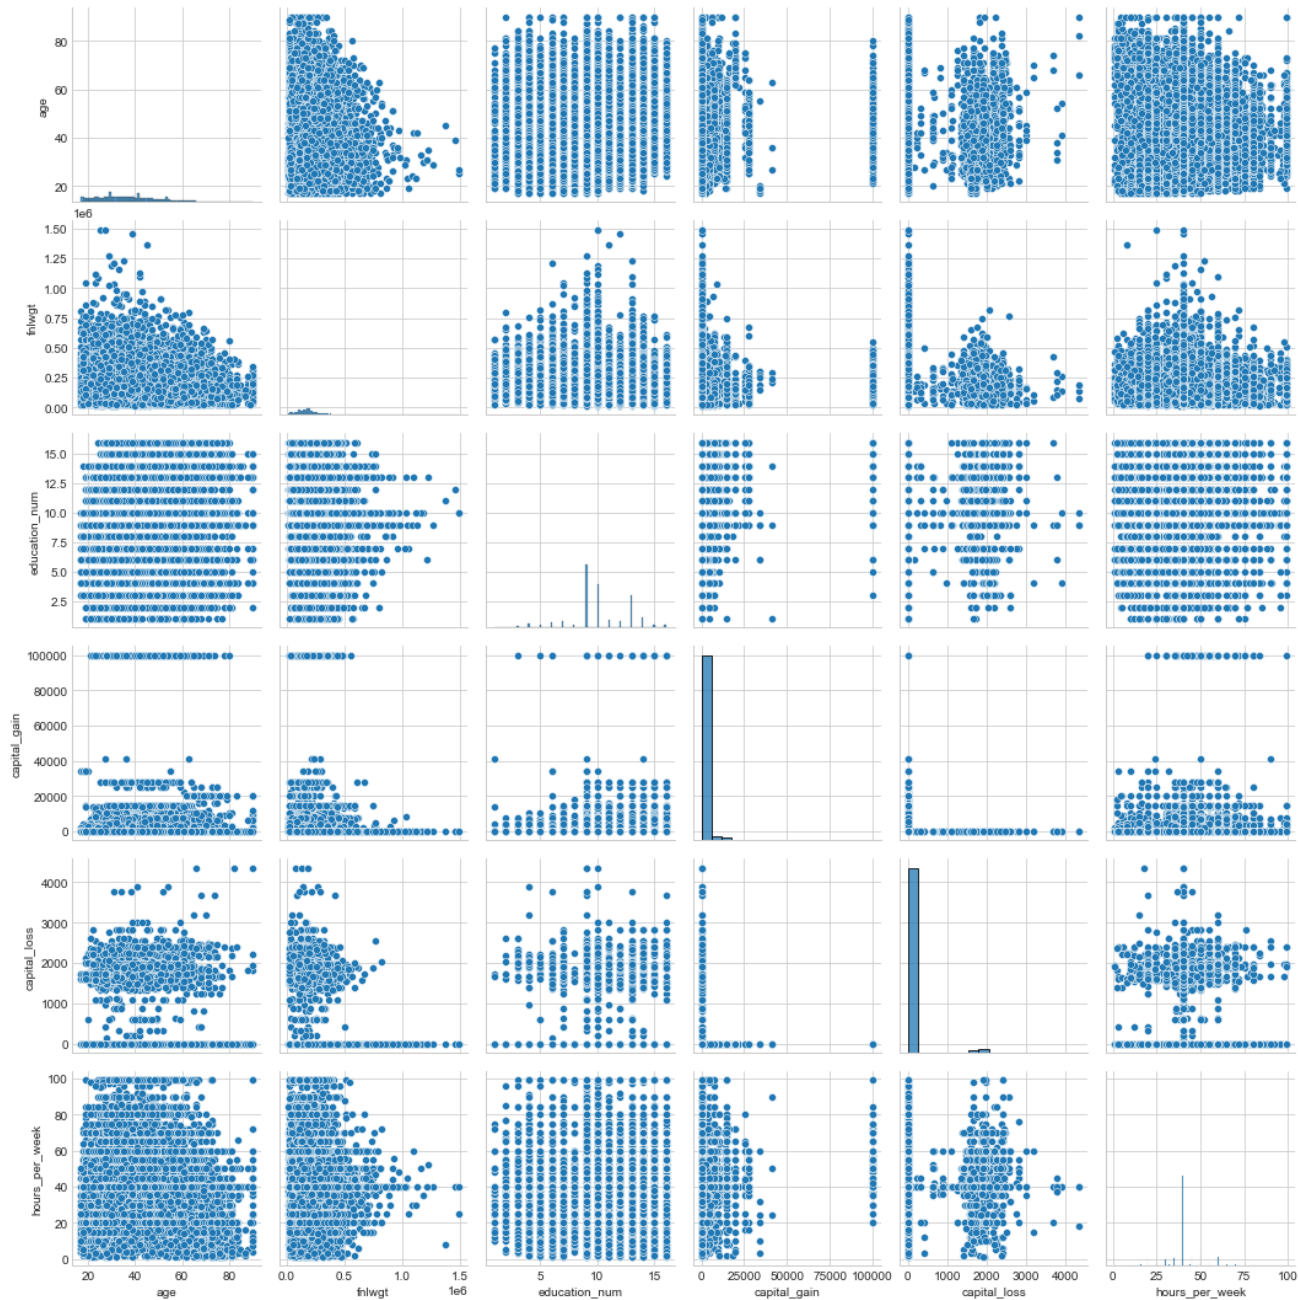
Multivariate Analysis

In [7]:

```
sns.pairplot(df)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x27987240e80>



Handling Duplicates

In [8]:

```
df.duplicated().sum()
```

Out[8]:

29

In [9]:

```
df = df.drop_duplicates()
```

In [10]:

```
df.duplicated().sum()
```

Out[10]:

0

Handling Null Values

In [11]:

```
df.isnull().sum()
```

Out[11]:

```
age                0
worktype          2799
fnlwgt            0
education         0
education_num     0
marital_status    0
occupation       2809
relationship      0
race             0
gender           0
capital_gain     0
capital_loss     0
hours_per_week   0
native_country   856
salary           0
dtype: int64
```

In [12]:

```
null_values = df.isnull().sum().sum()
if null_values == 0:
    print('No null values exist')
else:
    from sklearn.impute import SimpleImputer
    #imputing with most frequent values
    imputer = SimpleImputer(strategy='most_frequent', missing_values=np.nan)
    imputer = imputer.fit(df)
    df = pd.DataFrame(imputer.transform(df.loc[:, :]), columns = df.columns)
```

In [13]:

```
df.isnull().sum()
```

Out[13]:

```
age                0
worktype           0
fnlwgt            0
education         0
education_num     0
marital_status    0
occupation        0
relationship      0
race             0
gender           0
capital_gain     0
capital_loss     0
hours_per_week   0
native_country    0
salary           0
dtype: int64
```

Feature Encoding

Remember that Machine Learning algorithms perform Linear Algebra on Matrices, which means all features need have numeric values. The process of converting Categorical Features into values is called Encoding. Let's perform both One-Hot encoding.

Additional Resources: <http://pbpython.com/categorical-encoding.html> (<http://pbpython.com/categorical-encoding.html>)

In [14]:

```
# One Hot Encodes all Labels before Machine Learning
one_hot_cols = df.columns.tolist()
#one_hot_cols.remove('predclass')
dataset_bin_enc = pd.get_dummies(df, columns=one_hot_cols)

dataset_bin_enc.head()
```

Out[14]:

age_21	age_22	age_23	age_24	age_25	age_26	...	native_country_ Taiwan	native_country_ Thailand	native_country_ Trinidad&Tobago	native_country_ United-States	native_country_ Vietnam	native_countr Yugoslavia
0	0	0	0	0	0	...	0	0	0	1	0	
0	0	0	0	0	0	...	0	0	0	1	0	
0	0	0	0	0	0	...	0	0	0	1	0	
0	0	0	0	0	0	...	0	0	0	1	0	
0	0	0	0	0	0	...	0	0	0	0	0	





In [19]:

```
# Label Encode all Labels
df1 = df.apply(LabelEncoder().fit_transform)

df1.head()
```

Out[19]:

	age	worktype	fnlwgt	education	education_num	marital_status	occupation	relationship	race	gender	capital_gain	capital_loss	hours_per_week
0	22	6	3461	9	12	4	0	1	4	1	27	0	39
1	33	5	3788	9	12	2	3	0	4	1	0	0	12
2	21	3	18342	11	8	0	5	1	4	1	0	0	39
3	36	3	19995	1	6	2	5	0	2	1	0	0	39
4	11	3	25405	9	12	2	9	5	2	0	0	0	39

In [20]:

```
X = df1.drop('salary', axis = 1)
y = df1['salary']
```

In [23]:

```
X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=7,test_size=0.2)
```

In [31]:

```
'''Hyperparameters of Decision Tree Classifier'''
DTC_parameters = {
    'criterion' : ['gini', 'entropy', 'log_loss'],
    'splitter' : ['best', 'random'],
    'max_depth' : range(1,10,1),
    'min_samples_split' : range(2,10,2),
    'min_samples_leaf' : range(1,5,1),
    'max_features' : ['auto', 'sqrt', 'log2']
}

'''Hyperparameters of Bagging Classifier'''
Bagging_parameters = {
    'n_estimators' : [5, 10, 15],
    'max_samples' : range(2, 10, 1),
    'max_features' : range(2, 10, 3)
}

'''Hyperparameters of Random Forest Classifier'''
RFC_parameters = {
    'criterion' : ['gini', 'entropy', 'log_loss'],
    'max_depth' : range(1, 10, 1),
    'min_samples_split' : range(2, 10, 2),
    'min_samples_leaf' : range(1, 10, 1),
}

'''Hyperparameters of Random Forest Classifier'''
ETC_parameters = {
    'n_estimators' : [10,20,30],
    'criterion' : ['gini', 'entropy', 'log_loss'],
    'max_depth' : range(2,10,1),
    'min_samples_split' : range(2,10,2),
    'min_samples_leaf' : range(1,5,1),
    'max_features' : ['sqrt', 'log2']
}

'''Hard and Soft Voting Classifier'''
lr = LogisticRegression(multi_class='multinomial', random_state=7)
rfc = RandomForestClassifier(n_estimators=50, random_state=7)
svc = SVC(probability=True, random_state=7)

'''All Models'''
models = {
    1 : DecisionTreeClassifier(),
    2 : GridSearchCV(estimator = DecisionTreeClassifier(), param_grid = DTC_parameters, verbose=2, n_jobs = -1, cv=3),
    3 : BaggingClassifier(),
    4 : GridSearchCV(estimator = BaggingClassifier(), param_grid = Bagging_parameters, verbose=2, n_jobs = -1, cv=3),
    5 : RandomForestClassifier(),
    6 : GridSearchCV(estimator = RandomForestClassifier(), param_grid = RFC_parameters, verbose=2, n_jobs = -1, cv=3),
    7 : ExtraTreesClassifier(),
    8 : GridSearchCV(estimator = ExtraTreesClassifier(), param_grid = ETC_parameters, verbose=2, n_jobs = -1, cv=3),
    9 : VotingClassifier(estimators = [('lr', lr), ('rfc', rfc),('svc', svc)], voting='hard'),
    10 : VotingClassifier(estimators = [('lr', lr), ('rfc', rfc),('svc', svc)], voting='soft')
}
```

In [32]:

```
map_keys = list(models.keys())
```

In [33]:

```
# Get model name using id from linear_model_collection
def get_model_building_technique_name(num):
    if num == 1:
        return 'DecisionTreeClassifier'
    if num == 2:
        return 'GridSearchCV_DecisionTreeClassifier'
    if num ==3:
        return 'BaggingClassifier'
    if num == 4:
        return 'GridSearchCV_BaggingClassifier'
    if num == 5:
        return 'RandomForestClassifier'
    if num == 6:
        return 'GridSearchCV_RandomForestClassifier'
    if num == 7:
        return 'ExtraTreesClassifier'
    if num == 8:
        return 'GridSearchCV_ExtraTreesClassifier'
    if num == 9:
        return 'VotingClassifier_Hard'
    if num ==10:
        return 'VotingClassifier_Soft'
    return ''
```

In [36]:

```
results = [];  
for key_index in range(len(map_keys)):  
    key = map_keys[key_index]  
    if key in [1,2,3,4,5,6,7,8]:  
        model = models[key]  
        print(key)  
        model.fit(X_train, y_train)  
  
        '''Test Accuracy'''  
        y_pred = model.predict(X_test)  
  
        Accuracy_Test = accuracy_score(y_test, y_pred)  
        conf_mat_Test = confusion_matrix(y_test, y_pred)  
        true_positive_Test = conf_mat_Test[0][0]  
        false_positive_Test = conf_mat_Test[0][1]  
        false_negative_Test = conf_mat_Test[1][0]  
        true_negative_Test = conf_mat_Test[1][1]  
        Precision_Test = true_positive_Test / (true_positive_Test + false_positive_Test)  
        Recall_Test = true_positive_Test / (true_positive_Test + false_negative_Test)  
        F1_Score_Test = 2 * (Recall_Test * Precision_Test) / (Recall_Test + Precision_Test)  
        #AUC_Test = roc_auc_score(y_test, y_pred)  
  
        '''Train Accuracy'''  
        y_pred_train = model.predict(X_train)  
  
        Accuracy_Train = accuracy_score(y_train, y_pred_train)  
        conf_mat_Train = confusion_matrix(y_train, y_pred_train)  
        true_positive_Train = conf_mat_Train[0][0]  
        false_positive_Train = conf_mat_Train[0][1]  
        false_negative_Train = conf_mat_Train[1][0]  
        true_negative_Train = conf_mat_Train[1][1]  
        Precision_Train = true_positive_Train / (true_positive_Train + false_positive_Train)  
        Recall_Train = true_positive_Train / (true_positive_Train + false_negative_Train)  
        F1_Score_Train = 2 * (Recall_Train * Precision_Train) / (Recall_Train + Precision_Train)  
        #AUC_Train = roc_auc_score(y_train, y_pred_train)  
  
        results.append({  
            'Model Name' : get_model_building_technique_name(key),  
            'Trained Model' : model,  
            'Accuracy_Test' : Accuracy_Test,  
            'Precision_Test' : Precision_Test,  
            'Recall_Test' : Recall_Test,  
            'F1_Score_Test' : F1_Score_Test,  
            #'AUC_Test' : AUC_Test,  
            'Accuracy_Train' : Accuracy_Train,  
            'Precision_Train' : Precision_Train,  
            'Recall_Train' : Recall_Train,  
            'F1_Score_Train' : F1_Score_Train,  
            #'AUC_Train' : AUC_Train  
        })  
    else:  
        key = map_keys[key_index]  
        model = models[key]  
        print(key)  
        model.fit(X_train, y_train)  
  
        '''Test Accuracy'''  
        y_pred = model.predict(X_test)  
  
        Accuracy_Test = accuracy_score(y_test, y_pred)  
        conf_mat_Test = confusion_matrix(y_test, y_pred)  
        true_positive_Test = conf_mat_Test[0][0]  
        false_positive_Test = conf_mat_Test[0][1]  
        false_negative_Test = conf_mat_Test[1][0]  
        true_negative_Test = conf_mat_Test[1][1]  
        Precision_Test = true_positive_Test / (true_positive_Test + false_positive_Test)  
        Recall_Test = true_positive_Test / (true_positive_Test + false_negative_Test)  
        F1_Score_Test = 2 * (Recall_Test * Precision_Test) / (Recall_Test + Precision_Test)  
        #AUC_Test = roc_auc_score(y_test, y_pred)  
  
        '''Train Accuracy'''  
        y_pred_train = model.predict(X_train)  
  
        Accuracy_Train = accuracy_score(y_train, y_pred_train)  
        conf_mat_Train = confusion_matrix(y_train, y_pred_train)  
        true_positive_Train = conf_mat_Train[0][0]  
        false_positive_Train = conf_mat_Train[0][1]  
        false_negative_Train = conf_mat_Train[1][0]  
        true_negative_Train = conf_mat_Train[1][1]  
        Precision_Train = true_positive_Train / (true_positive_Train + false_positive_Train)  
        Recall_Train = true_positive_Train / (true_positive_Train + false_negative_Train)  
        F1_Score_Train = 2 * (Recall_Train * Precision_Train) / (Recall_Train + Precision_Train)  
        #AUC_Train = roc_auc_score(y_train, y_pred_train)  
  
        results.append({  
            'Model Name' : get_model_building_technique_name(key),  
            'Trained Model' : model,  
            'Accuracy_Test' : Accuracy_Test,  
            'Precision_Test' : Precision_Test,  
            'Recall_Test' : Recall_Test,  
            'F1_Score_Test' : F1_Score_Test,  
            #'AUC_Test' : AUC_Test,  
            'Accuracy_Train' : Accuracy_Train,  
            'Precision_Train' : Precision_Train,  
            'Recall_Train' : Recall_Train,  
            'F1_Score_Train' : F1_Score_Train,  
            #'AUC_Train' : AUC_Train  
        })
```

1  
2  
Fitting 3 folds for each of 2592 candidates, totalling 7776 fits  
3  
4  
Fitting 3 folds for each of 72 candidates, totalling 216 fits  
5  
6  
Fitting 3 folds for each of 972 candidates, totalling 2916 fits  
7  
8  
Fitting 3 folds for each of 2304 candidates, totalling 6912 fits  
9  
10

In [37]:

```
result_df = pd.DataFrame(results)
result_df
```

Out[37]:

	Model Name	Trained Model	Accuracy_Test	Precision_Test	Recall_Test	F1_Score_Test	Accu
0	DecisionTreeClassifier	DecisionTreeClassifier()	0.449862	0.647032	0.660094	0.653498	
1	GridSearchCV_DecisionTreeClassifier	GridSearchCV(cv=3, estimator=DecisionTreeClass...	0.562737	0.999149	0.662904	0.797014	
2	BaggingClassifier	(DecisionTreeClassifier(random_state=164281282...	0.512650	0.827063	0.661563	0.735113	
3	GridSearchCV_BaggingClassifier	GridSearchCV(cv=3, estimator=BaggingClassifier...	0.488989	1.000000	0.665820	0.799390	
4	RandomForestClassifier	(DecisionTreeClassifier(max_features='auto', r...	0.534160	0.884201	0.663588	0.758172	
5	GridSearchCV_RandomForestClassifier	GridSearchCV(cv=3, estimator=RandomForestClass...	0.570112	0.999788	0.662540	0.796954	
6	ExtraTreesClassifier	(ExtraTreeClassifier(random_state=1336411775),...	0.512855	0.828571	0.662358	0.736200	
7	GridSearchCV_ExtraTreesClassifier	GridSearchCV(cv=3, estimator=ExtraTreesClassif...	0.566322	1.000000	0.663687	0.797851	
8	VotingClassifier_Hard	VotingClassifier(estimators=[('lr',\n ...	0.544812	1.000000	0.664411	0.798374	
9	VotingClassifier_Soft	VotingClassifier(estimators=[('lr',\n ...	0.545631	0.995447	0.665100	0.797414	

In [ ]:

In [ ]: